

```
In [1]: import pandas as pd
import numpy as np
```

```
In [2]: import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [3]: import statsmodels.formula.api as smf
from sklearn.metrics import r2_score
```

```
In [4]: from sklearn.linear_model import LinearRegression
```

```
In [5]: from IPython.display import HTML
```

```
In [6]: import statsmodels.api as sm
```

```
/Users/kcarnold/anaconda3/envs/py36/lib/python3.6/site-packages/statsmodels/comp
at/pandas.py:56: FutureWarning: The pandas.core.datetools module is deprecated a
nd will be removed in a future version. Please use the pandas.tseries module ins
tead.
    from pandas.core import datetools
```


```
In [7]: from textrec.paths import paths
```

```
In [8]: dataset = pd.read_csv(paths.data / 'num_details_training_set.csv')
```

```
In [9]: sns.distplot(dataset.num_details)
```

```
/Users/kcarnold/anaconda3/envs/py36/lib/python3.6/site-packages/scipy/stats/stat
s.py:1706: FutureWarning: Using a non-tuple sequence for multidimensional indexi
ng is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future thi
s will be interpreted as an array index, `arr[np.array(seq)]`, which will result
either in an error or a different result.
    return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x1c1afe02e8>
```



Preprocessing

Strip off punctuation; it just throws off token counts and probs. (We get a few percent boost in r^2 because of this.)

```
In [10]: dataset['text'] = dataset.text.str.strip().str.rstrip('.')
```

```
In [11]: def strip_uninformative(text):
          text = text.strip()
          for beginning in ['there is', 'there are', 'a view of', 'a photo of', 'a photo
shows']:
              beginning = beginning + ' '
              if text.startswith(beginning):
                  text = text[len(beginning):]
              return strip_uninformative(text)
          return text
```

```
In [12]: strip_uninformative('there is a view of a red thing there')
```

```
Out[12]: 'a red thing there'
```

```
In [13]: dataset['text'] = dataset.text.apply(strip_uninformative)
```

Word Frequencies

```
In [14]: import wordfreq
```


```
In [15]: dataset['num_words'] = [
          len(wordfreq.tokenize(text, 'en'))
          for text in dataset.text]
```

```
In [16]: dataset['min_freq'] = [
          np.min([wordfreq.zipf_frequency(tok, 'en') for tok in wordfreq.tokenize(text,
          'en')])
          for text in dataset.text]
```

```
In [17]: dataset['mean_freq'] = [
          np.mean([wordfreq.zipf_frequency(tok, 'en') for tok in wordfreq.tokenize(text,
          'en')])
          for text in dataset.text]
```


```
In [18]: plt.scatter(dataset.min_freq, dataset.num_details)
```

```
Out[18]: <matplotlib.collections.PathCollection at 0x1c1de300f0>
```



```
In [19]: plt.scatter(dataset.mean_freq, dataset.num_details)
```

```
Out[19]: <matplotlib.collections.PathCollection at 0x1c1deae898>
```



Perplexity

The perplexity of a language model is a rough proxy for the amount of information that a text contains. The more details included, the more uncertainty the LM has; and redundant text doesn't get counted. It's not quite right for a few reasons:

- Typos, grammar errors, etc. also increase perplexity
- Unusual wording of the same concepts increases perplexity
- Using a word that's *more* common than expected increases perplexity.

But we'll try it anyway.

```
In [20]: from textrec import automated_analyses
         from textrec import onmt_model_2
```

```
/Users/kcarnold/anaconda3/envs/py36/lib/python3.6/site-packages/h5py/__init__.py
:36: FutureWarning: Conversion of the second argument of issubdtype from `float`
to `np.floating` is deprecated. In future, it will be treated as `np.float64 ==
np.dtype(float).type`.
```

```
from ._conv import register_converters as _register_converters
```

```
Loading ONMT models...
```

```
coco_lm_adam_acc_46.00_ppl_16.32_e10_nooptim.pt
```

```
Loading model parameters.
```

```
coco_cap_adam_acc_48.73_ppl_12.56_e10_nooptim.pt
```

```
Loading model parameters.
```

```
Ready.
```

```
Loading SpaCy...done
```

```
In [21]: automated_analyses.eval_logprobs_unconditional(dataset.text.iloc[0])
```

```
Out[21]: 3.3479643
```

```
In [22]: example_text = dataset.text.iloc[0]
         example_text
```

```
Out[22]: 'families stand around by the water flying kites on a sunny day'
```


```
In [23]: tokens = onmt_model_2.tokenize(example_text)
         logprobs = onmt_model_2.models['coco_lm'].eval_logprobs('.', tokens, use_eos=True)
         logprobs
```

```
Out[23]: array([1.0859766e+01, 4.1381788e+00, 2.3187706e+00, 7.0211720e+00,
                1.5096430e+00, 2.1149969e+00, 5.2354274e+00, 3.4142053e-01,
                5.1307883e+00, 1.5549884e+00, 1.9299134e+00, 3.4333759e-03,
                1.3650393e+00], dtype=float32)
```

```
In [24]: dataset['num_tokens'] = dataset.text.apply(lambda text: len(onmt_model_2.tokenize(
         text)))
         dataset['mean_logprob_uncond'] = dataset.text.apply(lambda text: automated_analyse
         s.eval_logprobs_unconditional(text))
         dataset['total_logprob_uncond'] = dataset.mean_logprob_uncond * (dataset.num_token
         s + 1)
```


```
In [27]: plt.scatter(dataset.num_tokens, dataset.num_details)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x1c432d8d30>
```



```
In [28]: plt.scatter(dataset.total_logprob_uncond, dataset.num_details)
```

```
Out[28]: <matplotlib.collections.PathCollection at 0x1c34f0da90>
```



Models

```
In [29]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 168 entries, 0 to 167
Data columns (total 9 columns):
image_id          168 non-null int64
text              168 non-null object
num_details       168 non-null int64
num_words         168 non-null int64
min_freq          168 non-null float64
mean_freq         168 non-null float64
num_tokens        168 non-null int64
mean_logprob_uncond 168 non-null float64
total_logprob_uncond 168 non-null float64
dtypes: float64(4), int64(4), object(1)
memory usage: 11.9+ KB
```

```
In [50]: dataset[dataset.num_words != dataset.num_tokens][['text', 'num_words', 'num_tokens']]
```

Out[50]:

	text	num_words	num_tokens
14	a man in a red shirt with two children are on a beach holding a multi-colored kite while other people fly kites in the background	26	25
19	a man in a red shirt is helping his children fly a large rainbow-colored kite	16	15
22	the image shows a railroad track with a train on it further out in the distance. multiple white buildings hug the side if the track, with some woo...	30	31
35	a train passing a few small buildings, perhaps the station	10	11
41	a landscape of a train stop with an old-looking brownish train and a few brightly colored buildings to one side	21	20
43	a black-and-white picture of a young couple cutting the wedding cake with the help of a young photographer at the wedding event	24	22
45	a woman is standing next to a couple in front of a cake with a knife in it and holding the other woman's hands	24	25
56	a husband, bride and female all stand in front of a table holding a knife cutting a cake	18	19
80	a man-woman gracefully riding a wave using a surfboard	10	9
81	a surfer is riding a wave the water looks so refreshing it's a beautiful day	15	16
86	a double-decker bus drives through a busy city street in london	12	11
89	a busy city street with cars, a large red bus and pedestrians going about their day	16	17
91	the photo shows a downtown scene of a city. there are old buildings everywhere, and a red bus is prominent in the middle of the road. many people ...	32	33
97	a busy city street with cars and people along the streets with high-rise buildings on both sides	18	17
99	a red double-decker bus driving down a street next to tall buildings and a cloudy sky in london	19	18
101	a red double-decker bus passes a group of people to its left while a black car looks to pass	20	19
111	a curious cat sits perched upon a table, next to a glass of wine	14	15
118	a brownish-orange cat with yellow eyes is look to his left past a glass of red wine	18	17
128	sliding glass, frosted, shower doors with a tan towel hanging on the handle and a white toilet with a blue floor rug	22	24
141	someone is using a shower but it's hard to see due to the opaque glass	15	16
147	a toilet paper sits on top of a toilet next to the sink, in a plain bathroom	17	18
153	a toilet has a roll of toilet paper on it, and there is a sink that matches it to the right	21	22
155	a sink, mirror and toilet, all in white with a roll of toilet paper on the toilet	17	19

```
In [30]: dataset.mean_freq.describe()
```

```
Out[30]: count    168.000000
mean      5.738019
std       0.283893
min       4.790000
25%      5.579911
50%      5.764444
75%      5.942990
max       6.323333
Name: mean_freq, dtype: float64
```

Let's try including the interaction of mean_freq and tokens. That's sorta like the total word frequency.. if we invert frequency to make `rarity` , then it's total rarity, or something proportional to unigram perplexity.

```
In [31]: dataset['mean_rarity'] = (7 - dataset.mean_freq) / 7
dataset['max_rarity'] = (7 - dataset.min_freq) / 7
dataset['total_rarity'] = dataset['mean_rarity'] * dataset['num_words']
```

```
In [32]: formulas = '''
C(image_id) + min_freq + mean_freq
C(image_id) + total_rarity
C(image_id) + num_tokens + total_rarity
C(image_id) + num_tokens + total_rarity + total_logprob_uncond + mean_logprob_uncond
C(image_id) + num_tokens + mean_rarity + max_rarity + total_rarity + total_logprob_uncond + mean_logprob_uncond + max_rarity*num_tokens
C(image_id) + num_tokens + mean_rarity + max_rarity + total_rarity + max_rarity*num_tokens
'''

models = {}
for formula in formulas:
    formula = formula.strip()
    if not formula:
        continue
    formula_full = 'num_details ~ ' + formula
    models[formula] = model = smf.ols(formula_full, dataset).fit()
    display(HTML(f'<h1>r^2={model.rsquared:.3f}: {formula}</h1>'))
    display(model.summary())
```


$r^2=0.385$: C(image_id) + min_freq + mean_freq

OLS Regression Results

Dep. Variable:	num_details	R-squared:	0.385			
Model:	OLS	Adj. R-squared:	0.350			
Method:	Least Squares	F-statistic:	10.98			
Date:	Wed, 10 Oct 2018	Prob (F-statistic):	3.40e-13			
Time:	09:49:30	Log-Likelihood:	-378.74			
No. Observations:	168	AIC:	777.5			
Df Residuals:	158	BIC:	808.7			
Df Model:	9					
Covariance Type:	nonrobust					
	coef	std err	t	P> t 	[0.025	0.975]
Intercept	3.2883	3.939	0.835	0.405	-4.492	11.068
C(image_id)[T.223777]	-2.3113	0.802	-2.882	0.005	-3.896	-0.727
C(image_id)[T.227326]	-2.1374	0.752	-2.841	0.005	-3.624	-0.651
C(image_id)[T.240275]	-4.3229	0.736	-5.877	0.000	-5.776	-2.870
C(image_id)[T.247576]	1.4757	0.754	1.956	0.052	-0.014	2.966
C(image_id)[T.275449]	-1.2033	0.749	-1.607	0.110	-2.682	0.276
C(image_id)[T.396295]	-1.1903	0.757	-1.572	0.118	-2.686	0.305
C(image_id)[T.431140]	0.1770	0.772	0.229	0.819	-1.349	1.703
min_freq	-1.5700	0.342	-4.595	0.000	-2.245	-0.895
mean_freq	1.4900	0.740	2.015	0.046	0.029	2.951
Omnibus:	17.294	Durbin-Watson:	2.051			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	23.230			
Skew:	0.628	Prob(JB):	9.03e-06			
Kurtosis:	4.320	Cond. No.	148.			

$r^2=0.713$: C(image_id) + total_rarity

OLS Regression Results

Dep. Variable:	num_details	R-squared:	0.713
Model:	OLS	Adj. R-squared:	0.699
Method:	Least Squares	F-statistic:	49.49
Date:	Wed, 10 Oct 2018	Prob (F-statistic):	2.32e-39
Time:	09:49:30	Log-Likelihood:	-314.55
No. Observations:	168	AIC:	647.1
Df Residuals:	159	BIC:	675.2
Df Model:	8		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.2352	0.473	4.722	0.000	1.300	3.170
C(image_id)[T.223777]	-3.6844	0.499	-7.381	0.000	-4.670	-2.698
C(image_id)[T.227326]	-3.6428	0.503	-7.243	0.000	-4.636	-2.649
C(image_id)[T.240275]	-3.6986	0.500	-7.399	0.000	-4.686	-2.711
C(image_id)[T.247576]	-0.1979	0.503	-0.393	0.695	-1.192	0.796
C(image_id)[T.275449]	-1.7292	0.499	-3.464	0.001	-2.715	-0.743
C(image_id)[T.396295]	-1.9313	0.499	-3.868	0.000	-2.917	-0.945
C(image_id)[T.431140]	-0.9063	0.500	-1.813	0.072	-1.893	0.081
total_rarity	1.9068	0.126	15.109	0.000	1.658	2.156

Omnibus:	1.608	Durbin-Watson:	1.877
Prob(Omnibus):	0.447	Jarque-Bera (JB):	1.227
Skew:	-0.183	Prob(JB):	0.541
Kurtosis:	3.204	Cond. No.	25.2

$r^2=0.736$: C(image_id) + num_tokens + total_rarity

OLS Regression Results

Dep. Variable:	num_details	R-squared:	0.736			
Model:	OLS	Adj. R-squared:	0.721			
Method:	Least Squares	F-statistic:	48.90			
Date:	Wed, 10 Oct 2018	Prob (F-statistic):	3.07e-41			
Time:	09:49:30	Log-Likelihood:	-307.73			
No. Observations:	168	AIC:	635.5			
Df Residuals:	158	BIC:	666.7			
Df Model:	9					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.0451	0.459	4.457	0.000	1.139	2.951
C(image_id)[T.223777]	-3.9423	0.486	-8.112	0.000	-4.902	-2.982
C(image_id)[T.227326]	-4.0243	0.496	-8.120	0.000	-5.003	-3.045
C(image_id)[T.240275]	-3.6838	0.482	-7.650	0.000	-4.635	-2.733
C(image_id)[T.247576]	-0.4408	0.489	-0.901	0.369	-1.407	0.526
C(image_id)[T.275449]	-1.8718	0.482	-3.880	0.000	-2.825	-0.919
C(image_id)[T.396295]	-1.9181	0.481	-3.988	0.000	-2.868	-0.968
C(image_id)[T.431140]	-1.3426	0.496	-2.707	0.008	-2.322	-0.363
num_tokens	0.1481	0.041	3.656	0.000	0.068	0.228
total_rarity	1.1980	0.229	5.236	0.000	0.746	1.650
Omnibus:	5.045	Durbin-Watson:	1.760			
Prob(Omnibus):	0.080	Jarque-Bera (JB):	5.034			
Skew:	-0.281	Prob(JB):	0.0807			
Kurtosis:	3.635	Cond. No.	140.			

$r^2=0.743$: C(image_id) + num_tokens + total_rarity + total_logprob_uncond + mean_logprob_uncond

OLS Regression Results

Dep. Variable:	num_details	R-squared:	0.743			
Model:	OLS	Adj. R-squared:	0.725			
Method:	Least Squares	F-statistic:	41.00			
Date:	Wed, 10 Oct 2018	Prob (F-statistic):	1.84e-40			
Time:	09:49:30	Log-Likelihood:	-305.41			
No. Observations:	168	AIC:	634.8			
Df Residuals:	156	BIC:	672.3			
Df Model:	11					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.7258	1.160	0.625	0.533	-1.566	3.018
C(image_id)[T.223777]	-3.7699	0.490	-7.699	0.000	-4.737	-2.803
C(image_id)[T.227326]	-4.0356	0.493	-8.182	0.000	-5.010	-3.061
C(image_id)[T.240275]	-3.6715	0.478	-7.674	0.000	-4.617	-2.726
C(image_id)[T.247576]	-0.4646	0.486	-0.956	0.340	-1.424	0.495
C(image_id)[T.275449]	-1.8630	0.480	-3.878	0.000	-2.812	-0.914
C(image_id)[T.396295]	-1.7238	0.492	-3.502	0.001	-2.696	-0.751
C(image_id)[T.431140]	-1.3638	0.493	-2.768	0.006	-2.337	-0.391
num_tokens	0.2491	0.081	3.091	0.002	0.090	0.408
total_rarity	1.3737	0.246	5.574	0.000	0.887	1.861
total_logprob_uncond	-0.0292	0.018	-1.637	0.104	-0.064	0.006
mean_logprob_uncond	0.2810	0.277	1.014	0.312	-0.266	0.828
Omnibus:	3.726	Durbin-Watson:	1.837			
Prob(Omnibus):	0.155	Jarque-Bera (JB):	3.774			
Skew:	-0.180	Prob(JB):	0.152			
Kurtosis:	3.640	Cond. No.	731.			

$r^2=0.745$: C(image_id) + num_tokens + mean_rarity + max_rarity + total_rarity + total_logprob_uncond + mean_logprob_uncond + max_rarity*num_tokens

OLS Regression Results

Dep. Variable:	num_details	R-squared:	0.745			
Model:	OLS	Adj. R-squared:	0.722			
Method:	Least Squares	F-statistic:	31.92			
Date:	Wed, 10 Oct 2018	Prob (F-statistic):	2.57e-38			
Time:	09:49:30	Log-Likelihood:	-304.77			
No. Observations:	168	AIC:	639.5			
Df Residuals:	153	BIC:	686.4			
Df Model:	14					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.8745	2.096	0.894	0.372	-2.266	6.015
C(image_id)[T.223777]	-3.7092	0.551	-6.738	0.000	-4.797	-2.622
C(image_id)[T.227326]	-4.0110	0.514	-7.803	0.000	-5.026	-2.996
C(image_id)[T.240275]	-3.6637	0.485	-7.551	0.000	-4.622	-2.705
C(image_id)[T.247576]	-0.4367	0.513	-0.852	0.396	-1.450	0.576
C(image_id)[T.275449]	-1.8771	0.496	-3.784	0.000	-2.857	-0.897
C(image_id)[T.396295]	-1.7201	0.520	-3.308	0.001	-2.748	-0.693
C(image_id)[T.431140]	-1.3342	0.516	-2.584	0.011	-2.354	-0.314
num_tokens	0.1609	0.126	1.274	0.204	-0.089	0.410
mean_rarity	-3.8772	8.361	-0.464	0.643	-20.394	12.640
max_rarity	-1.5791	3.463	-0.456	0.649	-8.421	5.263
total_rarity	1.5233	0.545	2.796	0.006	0.447	2.600
total_logprob_uncond	-0.0360	0.020	-1.797	0.074	-0.076	0.004
mean_logprob_uncond	0.3996	0.332	1.205	0.230	-0.256	1.055
max_rarity:num_tokens	0.1615	0.204	0.792	0.430	-0.241	0.564
Omnibus:	4.007	Durbin-Watson:	1.867			
Prob(Omnibus):	0.135	Jarque-Bera (JB):	4.201			
Skew:	-0.183	Prob(JB):	0.122			
Kurtosis:	3.683	Cond. No.	4.92e+03			

$r^2=0.737$: C(image_id) + num_tokens + mean_rarity + max_rarity + total_rarity + max_rarity*num_tokens

OLS Regression Results

Dep. Variable:	num_details	R-squared:	0.737			
Model:	OLS	Adj. R-squared:	0.716			
Method:	Least Squares	F-statistic:	36.13			
Date:	Wed, 10 Oct 2018	Prob (F-statistic):	7.70e-39			
Time:	09:49:30	Log-Likelihood:	-307.47			
No. Observations:	168	AIC:	640.9			
Df Residuals:	155	BIC:	681.5			
Df Model:	12					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	3.1419	2.030	1.548	0.124	-0.868	7.152
C(image_id)[T.223777]	-3.9445	0.545	-7.241	0.000	-5.021	-2.868
C(image_id)[T.227326]	-4.0215	0.519	-7.756	0.000	-5.046	-2.997
C(image_id)[T.240275]	-3.6782	0.489	-7.529	0.000	-4.643	-2.713
C(image_id)[T.247576]	-0.4441	0.518	-0.858	0.392	-1.467	0.578
C(image_id)[T.275449]	-1.9017	0.501	-3.799	0.000	-2.891	-0.913
C(image_id)[T.396295]	-1.9284	0.506	-3.813	0.000	-2.928	-0.929
C(image_id)[T.431140]	-1.3438	0.521	-2.578	0.011	-2.373	-0.314
num_tokens	0.0734	0.121	0.608	0.544	-0.165	0.312
mean_rarity	-1.7638	6.952	-0.254	0.800	-15.498	11.970
max_rarity	-1.4324	3.492	-0.410	0.682	-8.331	5.466
total_rarity	1.2691	0.492	2.581	0.011	0.298	2.241
max_rarity:num_tokens	0.1159	0.205	0.567	0.572	-0.288	0.520
Omnibus:	4.893	Durbin-Watson:	1.774			
Prob(Omnibus):	0.087	Jarque-Bera (JB):	4.818			
Skew:	-0.279	Prob(JB):	0.0899			
Kurtosis:	3.614	Cond. No.	1.08e+03			


Summary: num_details increases by 1.4 for each additional token of rarity.

Let's look at resids.

```
In [33]: model = models['C(image_id) + num_tokens + total_rarity']
```

```
In [34]: predicted = model.predict(dataset)
```

```
In [35]: plt.scatter(predicted, model.resid)
plt.axhline(0, color='r')
plt.xlabel('$\hat{y}$')
plt.ylabel('residual');
```



Ok, let's have a look at captions for which length and frequency don't predict num_details well.

```
In [36]: dsr = dataset.copy()
```

```
In [37]: dsr['resid'] = model.resid
dsr['resid_mag'] = model.resid.abs()
dsr['predicted'] = predicted
dsr.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 168 entries, 0 to 167
Data columns (total 15 columns):
image_id          168 non-null int64
text              168 non-null object
num_details       168 non-null int64
num_words         168 non-null int64
min_freq          168 non-null float64
mean_freq         168 non-null float64
num_tokens        168 non-null int64
mean_logprob_uncond 168 non-null float64
total_logprob_uncond 168 non-null float64
mean_rarity       168 non-null float64
max_rarity        168 non-null float64
total_rarity      168 non-null float64
resid             168 non-null float64
resid_mag         168 non-null float64
predicted         168 non-null float64
dtypes: float64(10), int64(4), object(1)
memory usage: 19.8+ KB
```

```
In [38]: pd.set_option('display.max_colwidth', 150)
```

```
In [39]: print("over-predicted:")
dsr[dsr.predicted.between(4,8)][['image_id text num_tokens resid predicted total_logprob_uncond num_details'].split()].sort_values('resid').iloc[:5]
```

over-predicted:

Out[39]:

	image_id	text	num_tokens	resid	predicted	total_logprob_uncond	num_details
3	200451	several multicolored kites with streamers are seen soaring above the heads of people	13	-3.965627	7.965627	54.609120	4
22	223777	the image shows a railroad track with a train on it further out in the distance. multiple white buildings hug the side if the track, with some woo...	31	-3.901685	7.901685	180.252899	4
9	200451	one kite flying over four other kites on a blue sky	11	-2.645860	6.645860	43.533666	4
103	247576	a double decker bus traveling down the middle of the street in the city streets	15	-2.493046	6.493046	36.642353	4
164	431140	toilet paper roll is on top of the toilet in a mellow yellow painted bathroom	15	-2.422958	6.422958	75.009583	4

```
In [40]: print("Under-predicted")
dsr[dsr.predicted.between(4,8)][['image_id text num_tokens resid predicted total_logprob_uncond num_details'].split()].sort_values('resid').iloc[-5:]
```

Under-predicted

Out[40]:

	image_id	text	num_tokens	resid	predicted	total_logprob_uncond	num_details
120	275449	a half full glass of red wine on a table in front of a calico cat	16	2.541694	5.458306	38.410653	8
145	396295	a tan towel is hanging from a chrome handle on a textured glass shower door	15	2.543200	6.456800	81.209793	9
2	200451	a man and his two children are flying multicolored kites on a sandy beach	14	2.646132	7.353868	39.750552	10
7	200451	a man flies a butterfly kite with his two daughters	10	3.061945	5.938055	37.889245	9
151	431140	a bathroom with a white sink and white toilet. a roll of unwrapped toilet paper sits on the bowl	19	3.144114	7.855886	97.487974	11

I notice:

- We can generally do surprisingly well on this task using total rarity. We can explain about 74% of the variance in details.
- Some of the over-predicts actually have more details than I gave them credit for. Some of the under-predicts are less detailed.
- Some of the over-predicted just have extra words ("there is" one kite; "a view of" a bathroom, "in the city streets"); I went back and stripped them off and the above reflects that. (we get a boost of about 0.01 R^2.)

Since this model has image only as a slope (should be random but alas I'm lazy), we can still get relative details measures.

Aside: random-effects model.

```
In [41]: md = smf.mixedlm("num_details ~ mean_rarity + num_tokens + total_rarity", dataset,
                        groups=dataset["image_id"])

md.fit().summary()
```

Out[41]:

Model:	MixedLM	Dependent Variable:	num_details
No. Observations:	168	Method:	REML
No. Groups:	8	Scale:	2.4408
Min. group size:	21	Likelihood:	-323.9069
Max. group size:	21	Converged:	Yes
Mean group size:	21.0		

	Coef.	Std.Err.	z	P> z	[0.025	0.975]
Intercept	0.371	1.431	0.259	0.795	-2.434	3.177
mean_rarity	-2.535	6.720	-0.377	0.706	-15.707	10.636
num_tokens	0.118	0.085	1.392	0.164	-0.048	0.285
total_rarity	1.360	0.465	2.924	0.003	0.449	2.272
groups RE	2.372	0.870				

```
In [42]: md = smf.mixedlm("num_details ~ total_rarity", dataset, groups=dataset["image_id"])
mdf = md.fit()
mdf.summary()
```

```
Out[42]:
```

Model:	MixedLM	Dependent Variable:	num_details
No. Observations:	168	Method:	REML
No. Groups:	8	Scale:	2.6164
Min. group size:	21	Likelihood:	-330.8710
Max. group size:	21	Converged:	Yes
Mean group size:	21.0		

	Coef.	Std.Err.	z	P> z	[0.025	0.975]
Intercept	0.251	0.641	0.392	0.695	-1.006	1.508
total_rarity	1.911	0.126	15.161	0.000	1.664	2.158
groups RE	2.295	0.817				

```
In [43]: r2_score(dataset['num_details'], mdf.predict(dataset))
```

```
Out[43]: 0.4685851105546911
```

I don't understand why that R² score is much smaller than the fixed-effects version. Probably we have different parameters.

```
In [44]: model.params
```

```
Out[44]:
```

Intercept	2.045098
C(image_id)[T.223777]	-3.942306
C(image_id)[T.227326]	-4.024281
C(image_id)[T.240275]	-3.683757
C(image_id)[T.247576]	-0.440750
C(image_id)[T.275449]	-1.871804
C(image_id)[T.396295]	-1.918077
C(image_id)[T.431140]	-1.342626
num_tokens	0.148146
total_rarity	1.198048
dtype:	float64

```
In [45]: mdf.params
```

```
Out[45]:
```

Intercept	0.251184
total_rarity	1.910701
groups RE	0.877341
dtype:	float64

```
In [46]: re_params = pd.Series({k: v.iloc[0] for k, v in mdf.random_effects.items()})
re_params
```

```
Out[46]: 200451    1.872634
223777    -1.622234
227326    -1.584410
240275    -1.634757
247576    1.683106
275449    0.232425
396295    0.041020
431140    1.012215
dtype: float64
```

```
In [47]: fixed_params = pd.Series({int(k[14:-1]): v for k, v in model.params.items() if k.startswith('C')})
fixed_params
```

```
Out[47]: 223777    -3.942306
227326    -4.024281
240275    -3.683757
247576    -0.440750
275449    -1.871804
396295    -1.918077
431140    -1.342626
dtype: float64
```

```
In [48]: d = pd.DataFrame(dict(fixed=fixed_params + model.params['Intercept'], random=re_params + mdf.params['Intercept']))
d['diff'] = d['fixed'] - d['random']
d['absdiff'] = d['fixed'].abs() - d['random'].abs()
d.mean(axis=0)
```

```
Out[48]: fixed    -0.415417
random      0.251184
diff        -0.399082
absdiff      0.008690
dtype: float64
```

The random-effects estimates are generally smaller.